

## Dataset Setup

In [ ]:

```
import pathlib, os, sys, operator, re, datetime
from functools import reduce
import numpy as np
import tensorflow as tf
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras import Model
import tensorflow_datasets as tfds
from tiny_imagenet import TinyImagenetDataset
from keras.models import load_model

# Enable or disable GPU
# To fully disable it, we need to hide all GPU devices from Tensorflow
# Make sure GPU is disabled for this inference part of the lab
ENABLE_GPU = True
# tf.debugging.set_log_device_placement(True)

if not ENABLE_GPU:
    tf.config.set_visible_devices([], 'GPU')

# Print Python and TF version, and where we are running
print(f'Running on Python Version: {sys.version}')
print(f'Using Tensorflow Version: {tf.__version__}')
if not tf.config.experimental.list_physical_devices("GPU"):
    print('Running on CPU')
else:
    print(f'Using GPU at: {tf.test.gpu_device_name()} (of {len(tf.config.expe
```

```
Running on Python Version: 3.6.8 (default, May 31 2023, 10:28:59)
[GCC 8.5.0 20210514 (Red Hat 8.5.0-18)]
Using Tensorflow Version: 2.6.2
Using GPU at: /device:GPU:0 (of 1 available)
```

```
In [ ]: # Original Source: https://github.com/ksachdeva/tiny-imagenet-tfds
# Class Version Source: https://github.com/MLuckydwyer/tiny-imagenet-tfds
# Setup our dataset
# -----

tiny_imagenet_builder = TinyImagenetDataset()

# this call (download_and_prepare) will trigger the download of the dataset
# and preparation (conversion to tfrecords)
#
# This will be done only once and on next usage tfds will
# use the cached version on your host.
tiny_imagenet_builder.download_and_prepare(download_dir="/tensorflow-dataset")

# class_names = tiny_imagenet_builder.info.features['label'].names
ds = tiny_imagenet_builder.as_dataset()
ds_train, ds_val = ds["train"], ds["validation"]
assert(isinstance(ds_train, tf.data.Dataset))
assert(isinstance(ds_val, tf.data.Dataset))

# Training Dataset
ds_train = ds_train.shuffle(1024).prefetch(tf.data.AUTOTUNE)

# Validation Dataset
ds_val = ds_val.shuffle(1024).prefetch(tf.data.AUTOTUNE)

# Dataset metadata
ds_info = tiny_imagenet_builder.info
```

## Working with the Dataset

```
In [ ]: # We need to read the "human readable" labels so we can translate with the nu
# Read the labels file (words.txt)
with open(os.path.abspath('wnids.txt'), 'r') as f:
    wnids = [x.strip() for x in f]

# Map wnids to integer labels
wnid_to_label = {wnid: i for i, wnid in enumerate(wnids)}
label_to_wnid = {v: k for k, v in wnid_to_label.items()}

# Use words.txt to get names for each class
with open(os.path.abspath('words.txt'), 'r') as f:
    wnid_to_words = dict(line.split('\t') for line in f)
    for wnid, words in wnid_to_words.items():
        wnid_to_words[wnid] = [w.strip() for w in words.split(',')]

class_names = [str(wnid_to_words[wnid]) for wnid in wnids]
```

```
In [ ]: # Helper function to get the label name
def img_class(img_data, idx=None):
    image, label, id, label_name = img_data["image"], img_data["label"], img_
    # Handle batches of images correctly
    if idx != None:
        image, label, id, label_name = img_data["image"][idx], img_data["labe

    return f"{label_name} (class index: {label} - id: {id})"

# Helper function to show basic info about an image
def img_info(img, idx=None, display=True, title_apend=""):
    image = img['image']

    # Print the class
    class_str = img_class(img, idx)
    print(f"Label: {class_str}")

    # Display the image
    if display:
        plt.figure()
        plt.title(title_apend + class_str)
        # Handle batches correctly
        if image.shape.ndims > 3:
            plt.imshow(image.numpy().reshape(64, 64, 3))
        else:
            plt.imshow(image.numpy())
```

```
In [ ]: # Print the dataset types and info
print("--- Train & Validation dataset info ---")
print(f"Train: {ds_train}")
print(f"Validation: {ds_val}")
print(f"Dataset Info: {ds_info}") # Uncomment to print Dataset info

print("\n--- Show an example image ---")
for example in ds_val.take(1):
    img_info(example)

print("\n Show some other examples")
tfds.show_examples(ds_val, ds_info, rows=3, cols=3)
```

--- Train & Validation dataset info ---

Train: <PrefetchDataset shapes: {id: (), image: (64, 64, 3), label: (), metadata: {label\_name: ()}}, types: {id: tf.string, image: tf.uint8, label: tf.int64, metadata: {label\_name: tf.string}}>

Validation: <PrefetchDataset shapes: {id: (), image: (64, 64, 3), label: (), metadata: {label\_name: ()}}, types: {id: tf.string, image: tf.uint8, label: tf.int64, metadata: {label\_name: tf.string}}>

Dataset Info: tfds.core.DatasetInfo(

name='tiny\_imagenet\_dataset',  
full\_name='tiny\_imagenet\_dataset/0.2.0',  
description="""

Tiny ImageNet Challenge is a similar challenge as ImageNet with a smaller dataset but

less image classes. It contains 200 image classes, a training


dataset of 100, 000 images, a validation dataset of 10, 000 images, and a test dataset of 10, 000 images. All images are of size 64x64.

```
"""
homepage='https://www.tensorflow.org/datasets/catalog/tiny_imagenet_dataset',
data_path='/home/tjfriedl/tensorflow_datasets/tiny_imagenet_dataset/0.2.0',
download_size=236.61 MiB,
dataset_size=215.40 MiB,
features=FeaturesDict({
    'id': Text(shape=(), dtype=tf.string),
    'image': Image(shape=(64, 64, 3), dtype=tf.uint8),
    'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=200),
    'metadata': FeaturesDict({
        'label_name': tf.string,
    }),
}),
supervised_keys=('image', 'label'),
disable_shuffling=False,
splits={
    'train': <SplitInfo num_examples=100000, num_shards=2>,
    'validation': <SplitInfo num_examples=10000, num_shards=1>,
},
citation="""@article{tiny-imagenet,
    author = {Li,Fei-Fei}, {Karpathy,Andrej} and {Johnson,Justin}""",
)
```

--- Show an example image ---

Label: b'dam, dike, dyke' (class index: 88 - id: b'n03160309')

b'dam, dike, dyke' (class index: 88 - id: b'n03160309')





103 (103)



193 (193)



119 (119)



39 (39)



125 (125)



73 (73)



125 (125)



144 (144)



68 (68)

Out[ ]:



103 (103)



193 (193)



119 (119)




39 (39)



125 (125)



73 (73)



```

In [ ]: # TODO: Print and visualize three inputs from the validation set
#       : Print the storage data type
#       : Print and note the dimensions of each image
#       : Print the memory required to store each image

# Sample Images
sample_imgs = []
for index, img_data in enumerate(ds_val.take(3)):
    sample_imgs.append(img_data)
    image, label, id, label_name = img_data["image"], img_data["label"], img_

    data_type = image.dtype # Supposed storage data type
    image_dimensions = image.shape # Supposed image shape
    dtype_size = image.dtype.size
    mem_per_element = image_dimensions.num_elements()
    memory_required = dtype_size * mem_per_element

    print(f'\n--- Image {index} ---')
    print(f"Storage Data Type: {data_type}") # Print out the storage data typ
    print(f"Image Dimensions: {image_dimensions}") # Print and note the dimen
    print(f"Memory required: {memory_required} (Bytes)") # Print the memory r
    img_info(img_data)

```

--- Image 0 ---

Storage Data Type: <dtype: 'uint8'>

Image Dimensions: (64, 64, 3)

Memory required: 12288 (Bytes)

Label: b'tarantula' (class index: 43 - id: b'n01774750')

--- Image 1 ---

Storage Data Type: <dtype: 'uint8'>

Image Dimensions: (64, 64, 3)

Memory required: 12288 (Bytes)

Label: b'meat loaf, meatloaf' (class index: 192 - id: b'n07871810')

--- Image 2 ---


Storage Data Type: <dtype: 'uint8'>

Image Dimensions: (64, 64, 3)

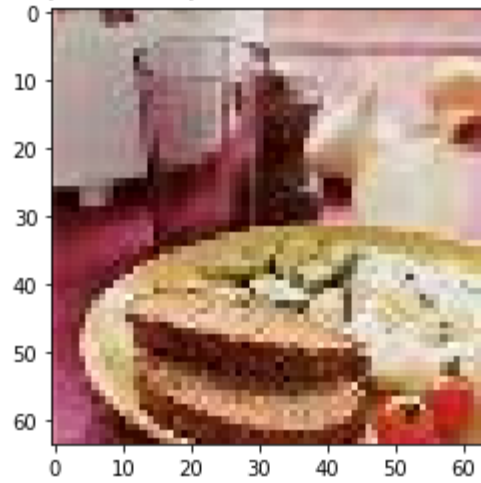
Memory required: 12288 (Bytes)

Label: b'sea slug, nudibranch' (class index: 165 - id: b'n01950731')

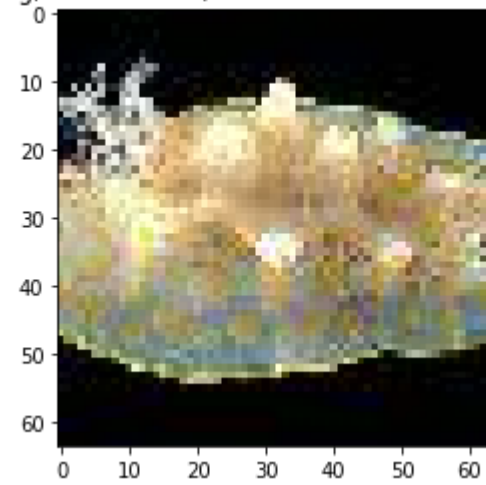
b'tarantula' (class index: 43 - id: b'n01774750')



b'meat loaf, meatloaf' (class index: 192 - id: b'n07871810')



b'sea slug, nudibranch' (class index: 165 - id: b'n01950731')





In [ ]:

```

# TODO: Export each of the three inputs to a binary file which will be used t
# NOTE: First flatten the array (ex: 4D --> 1D). So 64*64*3 = 12288 element 1

# Make a directory for our image data
img_dir = os.path.abspath('img_data')
pathlib.Path(img_dir).mkdir(exist_ok=True)

# Create a metadata file
metadata_file = open(os.path.join(img_dir, f'metadata.txt'), 'w')
metadata_file.write(f'Number\t\tDims\t\tClass Data\n')

# Export each image
for index, img_data in enumerate(sample_imgs):
    img_file = open(os.path.join(img_dir, f'image_{index}.bin'), 'wb')

    # TODO: Your Code Here
    image_data = img_data["image"].numpy() # Grabs the numerical image data f
    stringed_image_data = image_data.flatten() # Flattens the 4D array into b
    img_file.write(stringed_image_data.tobytes()) # Writes to the file in pro

    # print(stringed_image_data.shape) # Used to test we are acquiring 12288

    img_file.close()

    # Write the image metadata for reference later
    class_str = img_class(img_data)
    metadata_file.write(f'{index}\t\t\t{img_data["image"].shape}\t\t\t{class_str}')
metadata_file.close()

```

## Model Setup

In [ ]:

```

# TODO: Load the model
# Now we will load the H5 model! Please make sure the h5 model file is presen
# You can download this from the Canvas Page and place it in the same directo

# model_path = os.path.abspath("/home/dwyer/482/dev/CNN_TinyImageNet_2.h5")
model_path = os.path.abspath("CNN_TinyImageNet_2.h5")

# TODO: Your Code Here
model = tf.keras.models.load_model(model_path) # I believe this should do the

# TODO: Print a summary of the model
print(model.summary())
# NOTE: https://www.tensorflow.org/versions/r2.6/api_docs/python/tf/keras/Mod

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 60, 60, 32)	2432
conv2d_1 (Conv2D)	(None, 56, 56, 32)	25632
max_pooling2d (MaxPooling2D)	(None, 28, 28, 32)	0



**A1**  $\frac{1}{2}$   $\frac{1}{2}$   $\frac{1}{2}$

```
# Running inference on our model
# We can run an inference of our model by doing the following (we are doing ba
for example in ds_train.batch(1).take(1):
    img_info(example)

# Make a prediction
pred = model.predict(example["image"])
print(f'Raw 200 Class Weighted Prediction:\n{pred}') # Uncomment to see t

# What is out best guess?
best_guess = tf.math.argmax(pred, axis=1).numpy() # Our output is 200 wei
print(f'Best Guess [class index]: {class_names[best_guess[0]]} [{best_gue
print(f'Best Guess Confidence (percent / 1.0): {pred[0][best_guess]}')


# What are our top 15 guesses?
top_15 = tf.math.top_k(pred, k=15)
print(f'Top 15 Guesses (class index): {[f"{class_names[idx][0]} [{idx}]"
print(f'Top 15 Guesses Confidence (percent / 1.0): {top_15.values}')
```

9/11/23, 13:27

```

Top 15 Guesses (class index): ['[ 170]', '[ 0]', '[ 1]', '[ 2]', '[ 3]',
'[ 4]', '[ 5]', '[ 6]', '[ 7]', '[ 8]', '[ 9]', '[ 10]', '[ 11]', '[
[ 12]', '[ 13]']
Top 15 Guesses Confidence (percent / 1.0): [[1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
0. 0. 0. 0.]]
[b'bell pepper'] (class index: [188] - id: [b'n07720875'])

```



```
In [ ]: !who
```

```

In [ ]: # TODO: Run inference for our previous 3 sample images

# TODO: Your Code Here
for img_data in sample_imgs: # Start a for-loop in order to iterate through o
    img_info(img_data)

    prediction = model.predict(np.expand_dims(img_data["image"], axis=0))

    best_guess = tf.math.argmax(prediction, axis=1).numpy() # Our output is 2
    print(f'Best Guess [class index]: {class_names[best_guess[0]]} [{best_gue
    print(f'Best Guess Confidence (percent / 1.0): {prediction[0][best_guess]

    # What are our top 15 guesses?
    top_15 = tf.math.top_k(prediction, k=15)
    print(f'Top 15 Guesses (class index): {[f"{class_names[idx][0]} [{idx}]"
    print(f'Top 15 Guesses Confidence (percent / 1.0): {top_15.values}')

```


```

Label: b'tarantula' (class index: 43 - id: b'n01774750')
Best Guess [class index]: ['kimono'] [172]
Best Guess Confidence (percent / 1.0): [0.9999237]
Top 15 Guesses (class index): ['[ 172]', '[ 45]', '[ 173]', '[ 148]', '[
[ 10]', '[ 0]', '[ 1]', '[ 2]', '[ 3]', '[ 4]', '[ 5]', '[ 6]', '[
[ 7]', '[ 8]', '[ 9]']
Top 15 Guesses Confidence (percent / 1.0): [[9.9992371e-01 7.6236203e-05 1.19
01981e-19 8.2278266e-21 3.6131815e-36
0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00
0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00]]
Label: b'meat loaf, meatloaf' (class index: 192 - id: b'n07871810')
Best Guess [class index]: ['Christmas stocking'] [123]
Best Guess Confidence (percent / 1.0): [1.]
Top 15 Guesses (class index): ['[ 123]', '[ 63]', '[ 134]', '[ 61]', '[


```

```
[173]', '[ [45]', '[ [0]', '[ [1]', '[ [2]', '[ [3]', '[ [4]', '[ [5]', '[ [6]', '[ [7]', '[ [8]']
Top 15 Guesses Confidence (percent / 1.0): [[1.0000000e+00 1.5071672e-09 2.9508073e-19 5.7862415e-20 7.9649574e-35 1.9659285e-37 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00]]
Label: b'sea slug, nudibranch' (class index: 165 - id: b'n01950731')
Best Guess [class index]: ['comic book'] [173]
Best Guess Confidence (percent / 1.0): [0.9515173]
Top 15 Guesses (class index): ['[ [173]', '[ [61]', '[ [63]', '[ [46]', '[ [134]', '[ [146]', '[ [143]', '[ [0]', '[ [1]', '[ [2]', '[ [3]', '[ [4]', '[ [5]', '[ [6]', '[ [7]']
Top 15 Guesses Confidence (percent / 1.0): [[9.5151728e-01 4.8462652e-02 2.0052566e-05 1.9794566e-25 6.7907142e-27 9.8127139e-33 4.6507435e-35 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00 0.0000000e+00]]
```

b'tarantula' (class index: 43 - id: b'n01774750')



b'meat loaf, meatloaf' (class index: 192 - id: b'n07871810')



b'sea slug, nudibranch' (class index: 165 - id: b'n01950731')



In [ ]:

```
# TODO: Calculate the Top-1, Top-5, and Top-10 Accuracy of the validation data
total = acc_top1 = acc_top5 = acc_top10 = 0

# TODO: Your Code Here
for img_data in ds_val.batch(32): # Just creating another for loop in order to
    prediction = model.predict(img_data["image"].numpy())

    top_1 = tf.math.top_k(prediction, k=1).indices
    top_5 = tf.math.top_k(prediction, k=5).indices
    top_10 = tf.math.top_k(prediction, k=10).indices

    for i, img_label in enumerate(img_data['label'].numpy()):
        if img_label in top_1[i]:
            acc_top1 += 1
        if img_label in top_5[i]:
            acc_top5 += 1
        if img_label in top_10[i]:
            acc_top10 += 1

    total += 32

acc_top10 /= (total / 100) # Dividing by 1/100 multiplies :D
acc_top5 /= (total / 100)
acc_top1 /= (total / 100)

print("Accuracy (Top 10):", acc_top10, ", Accuracy (Top 5): ", acc_top5, ", Ac
```

Accuracy (Top 10): 15.225638977635784 ,Accuracy (Top 5): 12.649760383386582  
 , Accuracy (Top 1): 6.968849840255591

In [ ]:

```

# TODO: Print all of the possible classes of the dataset

train_classes = val_classes = 0

# TODO: Your Code Here
for image in ds_val:
    #class_name = image['metadata']['label_name'].numpy()
    val_classes += 1

# COMMENTED OUT FOR SUBMISSION

for image in ds_train:
    # class_name = image['metadata']['label_name'].numpy()
    train_classes += 1

print(class_names)
print("NUMBER OF CLASSES IN VAL DATASET: ", val_classes)
print("NUMBER OF CLASSES IN TRAIN DATASET", train_classes)

```

```

["['Egyptian cat']", "['reel']", "['volleyball']", "['rocking chair', 'rocker']", "['lemon']", "['bullfrog', 'Rana catesbeiana']", "['basketball']", "['cliff', 'drop', 'drop-off']", "['espresso']", "['\plunger\'], 'plumber's helper'", "['parking meter']", "['German shepherd', 'German shepherd dog', 'German police dog', 'alsatian']", "['dining table', 'board']", "['monarch', 'monarch butterfly', 'milkweed butterfly', 'Danaus plexippus']", "['brown bear', 'brown bear', 'Ursus arctos']", "['school bus']", "['pizza', 'pizza pie']", "['guinea pig', 'Cavia cobaya']", "['umbrella']", "['organ', 'pipe organ']", "['oboe', 'hautboy', 'hautbois']", "['maypole']", "['goldfish', 'Carassius auratus']", "['potpie']", "['hourglass']", "['seashore', 'coast', 'seacoast', 'sea-coast']", "['computer keyboard', 'keypad']", "['Arabian camel', 'dromedary', 'Camelus dromedarius']", "['ice cream', 'icecream']", "['nail']", "['space heater']", "['cardigan']", "['baboon']", "['snail']", "['coral reef']", "['albatross', 'mollymawk']", "['\spider web\'], 'spider's web'", "['sea cucumber', 'holothurian']", "['backpack', 'back pack', 'knapsack', 'packsack', 'rucksack', 'haversack']", "['Labrador retriever']", "['pretzel']", "['king penguin', 'Aptenodytes patagonica']", "['sulphur butterfly', 'sulfur butterfly']", "['tarantula']", "['lesser panda', 'red panda', 'panda', 'bear cat', 'cat bear', 'Ailurus fulgens']", "['pop bottle', 'soda bottle']", "['banana']", "['sock']", "['cockroach', 'roach']", "['projectile', 'missile']", "['beer bottle']", "['mantis', 'mantid']", "['freight car']", "['guacamole']", "['remote control', 'remote']", "['European fire salamander', 'Salamandra salamandra']", "['lakeside', 'lakeshore']", "['chimpanzee', 'chimp', 'Pan troglodytes']", "['pay-phone', 'pay-station']", "['fur coat']", "['alp']", "['lampshade', 'lamp shade']", "['torch']", "['abacus']", "['moving van']", "['barrel', 'cask']", "['tabby', 'tabby cat']", "['goose']", "['koala', 'koala bear', 'kangaroo bear', 'native bear', 'Phascolarctos cinereus']", "['bullet train', 'bullet']", "['CD player']", "['teapot']", "['birdhouse']", "['gazelle']", "['\academic gown\'], '\academic robe\'], 'judge's robe'", "['tractor']", "['ladybug', 'ladybeetle', 'lady beetle', 'ladybird', 'ladybird beetle']", "['miniskirt', 'mini']", "['golden retriever']", "['triumphal arch']", "['cannon']", "['neck brace']", "['sombrero']", "['gasmask', 'respirator', 'gas helmet']", "['candle', 'taper', 'wax light']", "['desk']", "['frying pan', 'frypan', 'skillet']", "['bee']", "['dam', 'dike', 'dyke']", "['spiny lobster', 'langouste', 'rock lobster', 'crawfish', 'crayfish', 'sea crawfish']", "['police van', 'police wagon', 'paddy wagon', 'patrol wagon', 'wagon', 'black Maria']", "['iPod']", "['punching bag', 'punch bag', 'punching ball', 'punchball']", "['beacon', 'lighthouse', 'beacon light', 'pharos']", "['jellyfish']", "['wok']", "['potter's wheel']", "['sandal']", "['pill bottle']", "['butcher shop', 'meat market']"]

```

```
arket']", "['slug']", "['hog', 'pig', 'grunter', 'squealer', 'Sus scrofa']]",
"['cougar', 'puma', 'catamount', 'mountain lion', 'painter', 'panther', 'Feli
s concolor']]", "['crane']]", "['vestment']]", "['\\dragonfly\\', '\\darning needl
e\\', 'devil\\'s darning needle', '\\sewing needle\\', '\\snake feeder\\', '\\snake
doctor\\', '\\mosquito hawk\\', '\\skeeter hawk\\']]", "['cash machine', 'cash disp
enser', 'automated teller machine', 'automatic teller machine', 'automated te
ller', 'automatic teller', 'ATM']]", "['mushroom']]", "['jinrikisha', 'ricksha
', 'rickshaw']]", "['water tower']]", "['chest']]", "['snorkel']]", "['sunglasses
', 'dark glasses', 'shades']]", "['fly']]", "['limousine', 'limo']]", "['black s
tork', 'Ciconia nigra']]", "['dugong', 'Dugong dugon']]", "['sports car', 'spor
t car']]", "['water jug']]", "['suspension bridge']]", "['ox']]", "['ice lolly',
'lolly', 'lollipop', 'popsicle']]", "['turnstile']]", "['Christmas stocking']]",
"['broom']]", "['scorpion']]", "['wooden spoon']]", "['picket fence', 'paling
']]", "['rugby ball']]", "['sewing machine']]", "['steel arch bridge']]", "['Pers
ian cat']]", "['refrigerator', 'icebox']]", "['barn']]", "['apron']]", "['Yorkshi
re terrier']]", "['swimming trunks', 'bathing trunks']]", "['stopwatch', 'stop
watch']]", "['lawn mower', 'mower']]", "['thatch', 'thatched roof']]", "['founta
in']]", "['black widow', 'Latrodectus mactans']]", "['bikini', 'two-piece']]",
"['plate']]", "['teddy', 'teddy bear']]", "['barbershop']]", "['confectionery',
'confectionary', 'candy store']]", "['beach wagon', 'station wagon', 'wagon',
'estate car', 'beach waggon', 'station waggon', 'waggon']]", "['scoreboard']]",
"['orange']]", "['flagpole', 'flagstaff']]", "['American lobster', 'Northern lo
bster', 'Maine lobster', 'Homarus americanus']]", "['trolleybus', 'trolley coa
ch', 'trackless trolley']]", "['drumstick']]", "['dumbbell']]", "['brass', 'memo
rial tablet', 'plaque']]", "['bow tie', 'bow-tie', 'bowtie']]", "['convertible
']]", "['bighorn', 'bighorn sheep', 'cimarron', 'Rocky Mountain bighorn', 'Roc
ky Mountain sheep', 'Ovis canadensis']]", "['orangutan', 'orang', 'orangutang
', 'Pongo pygmaeus']]", "['American alligator', 'Alligator mississippiensis']]",
"['centipede']]", "['syringe']]", "['go-kart']]", "['brain coral']]", "['sea slug
', 'nudibranch']]", "['cliff dwelling']]", "['mashed potato']]", "['viaduct']]",
"['military uniform']]", "['pomegranate']]", "['chain']]", "['kimono']]", "['comi
c book']]", "['trilobite']]", "['bison']]", "['pole']]", "['boa constrictor', 'Co
nstrictor constrictor']]", "['poncho']]", "['bathtub', 'bathing tub', 'bath', '
tub']]", "['grasshopper', 'hopper']]", "['walking stick', 'walkingstick', 'stic
k insect']]", "['Chihuahua']]", "['tailed frog', 'bell toad', 'ribbed toad', 't
ailed toad', 'Ascaphus trui']]", "['lion', 'king of beasts', 'Panthera leo']]",
"['altar']]", "['obelisk']]", "['beaker']]", "['bell pepper']]", "['bannister', '
banister', 'balustrade', 'balusters', 'handrail']]", "['bucket', 'pail']]", "['
magnetic compass']]", "['meat loaf', 'meatloaf']]", "['gondola']]", "['standard
poodle']]", "['acorn']]", "['lifeboat']]", "['binoculars', 'field glasses', 'ope
ra glasses']]", "['cauliflower']]", "['African elephant', 'Loxodonta africana
']"]]
```

NUMBER OF CLASSES IN VAL DATASET: 10000

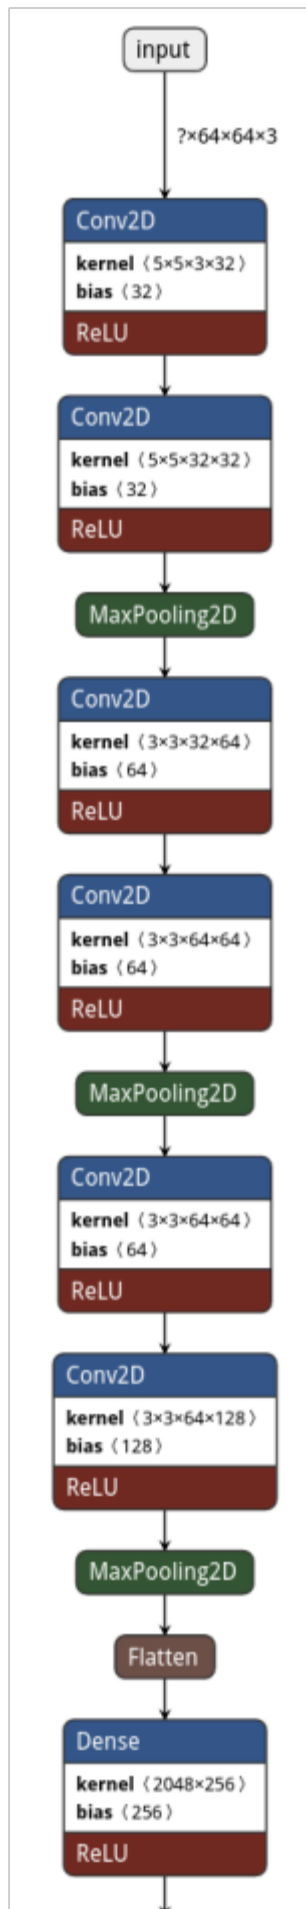
NUMBER OF CLASSES IN TRAIN DATASET 100000

## Model Exploration


In [ ]:

```
# TODO: Visualize the model in Netron and include an image here.
tf.keras.utils.plot_model(model, "model.png", show_shapes=True, show_dtype=True)

#My code for printing out the image created in Netron, using pyplot
plt.figure(figsize=(20, 20))
netron_image = plt.imread("netron_model.png") # Grabs saved image
plt.imshow(netron_image) # Displays the image
plt.axis('off') # Disables axis formatting in plot
plt.show() # Displays image
```







In [ ]:

```

# We can view the layer weights as well. Here we are pretending they are images
# TODO: Visualize the 2 convolutional layers filter sets (weights) (one at a time)

# TODO: Your Code Here

# Begin with our beginning input layer
# beginning_layer = model.layers[3]
# filter_weight, filter_biases = beginning_layer.get_weights()

# print(f'Beginning Layer: {beginning_layer.name}') # Prove name of layer
# print(f'Data Type: {filter_weight.dtype}') # Print dtype
# print(f'Dimensions: {filter_weight.ndim} total {filter_weight.shape}') # Print shape
# print(f'Required memory: {filter_weight.nbytes}B ({filter_weight.nbytes/1000000}MB)')

# filter_min, filter_max = filter_weight.min(), filter_weight.max()
# filter_weight = (filter_weight - filter_min) / (filter_max - filter_min)
# n_filters, i = 14, 1

# for foo in range(1, n_filters):
#     f = filter_weight[:, :, :, foo]
#     for j in range(3):
#         ax = plt.subplot(n_filters, 5, foo)
#         ax.set_xticks([])
#         ax.set_yticks([])

#         plt.imshow(f[:, :, j])
#         foo += 1
# plt.figure(figsize=(160, 160), dpi=10).show()

# # End with our later layer
# later_layer = model.layers[7]
# filter_weight, filter_biases = later_layer.get_weights()

# print(f'Ending Layer: {later_layer.name}') # Prove name of layer
# print(f'Data Type: {filter_weight.dtype}') # Print dtype
# print(f'Dimensions: {filter_weight.ndim} total {filter_weight.shape}') # Print shape
# print(f'Required memory: {filter_weight.nbytes}B ({filter_weight.nbytes/1000000}MB)')

# filter_min, filter_max = filter_weight.min(), filter_weight.max()
# later_layer = (filter_weight - filter_min) / (filter_max - filter_min)
# n_filters, i = 22, 1

# for foo in range(1, n_filters):
#     f = filter_weight[:, :, :, foo]
#     for j in range(5):
#         ax = plt.subplot(n_filters, 5, foo)
#         ax.set_xticks([])
#         ax.set_yticks([])

```

```

#         plt.imshow(f[:, :, j])
#         foo += 1
# plt.figure(figsize=(160, 160), dpi=10).show()

filters, biases = model.layers[0].get_weights()
f_min, f_max = filters.min(), filters.max()
filters = (filters - f_min) / (f_max - f_min)

n_filters, ix = 5, 1
for i in range(n_filters):
    f = filters[:, :, :, i]
    for j in range(3):
        ax = plt.subplot(n_filters, 3, ix)
        ax.set_xticks([])
        ax.set_yticks([])
        plt.imshow(f[:, :, j])
        plt.title(f'i={i}, j={j}')
        ix += 1
plt.show()

print(f'Beginning Layer: {model.layers[0].name}') # Prove name of layer
print(f'Data Type: {filters.dtype}') # Print dtype
print(f'Dimensions: {filters.ndim} total {filters.shape}') # Print #dims, + s
print(f'Required memory: {filters.nbytes}B ({filters.nbytes/1000}KB)') # Prin

filters, biases = model.layers[7].get_weights()
f_min, f_max = filters.min(), filters.max()
filters = (filters - f_min) / (f_max - f_min)

n_filters, ix = 5, 1
for i in range(n_filters):
    f = filters[:, :, :, i]
    for j in range(3):
        ax = plt.subplot(n_filters, 3, ix)
        ax.set_xticks([])
        ax.set_yticks([])
        plt.imshow(f[:, :, j])
        plt.title(f'i={i}, j={j}')
        ix += 1
plt.show()

print(f'Beginning Layer: {model.layers[7].name}') # Prove name of layer
print(f'Data Type: {filters.dtype}') # Print dtype
print(f'Dimensions: {filters.ndim} total {filters.shape}') # Print #dims, + s
print(f'Required memory: {filters.nbytes}B ({filters.nbytes/1000}KB)') # Prin

```

i=2, j=0



i=2, j=1



i=2, j=2



i=3, j=0



i=3, j=1




i=3, j=2





Beginning Layer: conv2d  
Data Type: float32  
Dimensions: 4 total (5, 5, 3, 32)  
Required memory: 9600B (9.6KB)



Beginning Layer: conv2d 5

In [ ]:

```

# We can again view the layer outputs as well. Here we are pretending they are
# TODO: Visualize the 2 convolutional layers outputs (intermediate feature maps)

# TODO: Your Code Here
# test_model = tf.keras.applications.vgg16.VGG16()

# for i in range(len(model.layers)):
#     layer = model.layers[i]
#     if 'conv' not in layer.name:
#         continue
#     print(i, layer.name, layer.output.shape)

# Run inference to get feature maps on earlier layer
sub_model = Model(inputs=model.inputs, outputs=model.layers[0].output)

#print(f"Printing feature maps for layer {model.layers[0].name}...")

image = sample_imgs[0]['image']
image = tf.keras.preprocessing.image.img_to_array(image)
image = np.expand_dims(image, axis=0)
image = tf.keras.applications.vgg16.preprocess_input(image)

feature_maps = sub_model.predict(image)
idx = 1
fig = plt.figure(figsize=(20, 20))
for i in range(1, feature_maps.shape[3]+1):

    if idx == 65:
        break

    plt.subplot(8, 8, i)
    plt.imshow(feature_maps[0, :, :, i-1])
    idx += 1

plt.show()

print(f'Beginning Layer: {model.layers[0].name}') # Prove name of layer
print(f'Data Type: {feature_maps.dtype}') # Print dtype
print(f'Dimensions: {feature_maps.ndim} total {feature_maps.shape}') # Print
print(f'Required memory: {feature_maps.nbytes}B ({feature_maps.nbytes/1000}KB)

# Run inference to get feature maps on end conv layer

sub_model = Model(inputs=model.input, outputs=model.layers[7].output)

#print(f"Printing feature maps for layer {model.layers[7].name}...")

image = sample_imgs[0]['image']
image = tf.keras.preprocessing.image.img_to_array(image)
image = np.expand_dims(image, axis=0)
image = tf.keras.applications.vgg16.preprocess_input(image)

feature_maps = sub_model.predict(image)

idx = 1
fig = plt.figure(figsize=(20, 20))

```

```

for i in range (1, feature_maps.shape[3]+1):


    if idx == 65:
        break

    plt.subplot(8, 8, i)
    plt.imshow(feature_maps[0, :, :, i-1])
    idx += 1


plt.show()

print(f'End Layer: {model.layers[7].name}') # Prove name of layer
print(f'Data Type: {feature_maps.dtype}') # Print dtype
print(f'Dimensions: {feature_maps.ndim} total {feature_maps.shape}') # Print
print(f'Required memory: {feature_maps.nbytes}B ({feature_maps.nbytes/1000}KB)

```



Beginning Layer: conv2d  
 Data Type: float32  
 Dimensions: 4 total (1, 60, 60, 32)  
 Required memory: 460800B (460.8KB)



In [ ]:

```

# TODO: Export the filters/weights so we can use them later
# Make a directory for our image data
model_dir = os.path.abspath('model_data')
pathlib.Path(model_dir).mkdir(exist_ok=True)

# Export each image
conv_index = dense_index = 1 # layer index starts from one
for layer_idx, layer in enumerate(model.layers):
    if re.match("conv", layer.name):
        weight_file_name = os.path.join(model_dir, f'conv{conv_index}_weights')
        bias_file_name = os.path.join(model_dir, f'conv{conv_index}_bias.bin')
        conv_index += 1
    elif re.match("dense", layer.name):
        weight_file_name = os.path.join(model_dir, f'dense{dense_index}_weights')
        bias_file_name = os.path.join(model_dir, f'dense{dense_index}_bias.bin')
        dense_index += 1
    else:
        continue
    # INPUT CODE BELOW
    weights, biases = layer.get_weights()
    weight_data, bias_data = weights.flatten(), biases.flatten()
    weight_file = open(weight_file_name, 'wb').write(weight_data.tobytes())
    bias_file = open(bias_file_name, 'wb').write(bias_data.tobytes())

print(f"All the convolution and dense (fully connected) weights and biases su

```

All the convolution and dense (fully connected) weights and biases successfully exported to input folders in /home/tjfriedl/Desktop/cpre\_487/lab1/model\_data directory

```
In [ ]: # TODO: Export the intermediate layer outputs for each of the input for all o
img_dir = os.path.abspath('img_data')
pathlib.Path(img_dir).mkdir(exist_ok=True)

for img_idx, img in enumerate(sample_imgs):
    file_dir = os.path.join(img_dir, f'test_input_{img_idx}')
    pathlib.Path(file_dir).mkdir(exist_ok=True)

    for layer_idx, layer in enumerate(model.layers):
        aux_model = tf.keras.Model(inputs=model.inputs, outputs=[layer.output

        # Store the intermediate output

        # TODO: Your Code Here
        aux_image = tf.keras.preprocessing.image.img_to_array(img['image'])
        aux_image = np.expand_dims(aux_image, axis=0)
        aux_image = tf.keras.applications.vgg16.preprocess_input(aux_image)

        feature_maps = aux_model.predict(aux_image)

        fmap_data = feature_maps.flatten()
        fmap_file_name = os.path.join(file_dir, f'layer_output_{layer_idx}.bi
        fmap_file = open(fmap_file_name, 'wb').write(fmap_data.tobytes())

print(f"All the corresponding intermediate layer outputs successfully exporte
```

All the corresponding intermediate layer outputs successfully exported to each input folder in the /home/tjfriedl/Desktop/cpre\_487/lab1/img\_data directory

## Tensorboard

```
In [ ]: # Setup for profiling
tf.profiler.experimental.ProfilerOptions(
    host_tracer_level=1, python_tracer_level=0, device_tracer_level=1
)

log_dir = os.path.abspath(os.path.join('log_data', datetime.datetime.now().st
pathlib.Path(log_dir).mkdir(exist_ok=True, parents=True)
```



In [ ]:

```

# TODO: Sample Profiling - Inference for a single image:

# Perform the inference profiling:

# THE ABOVE EXAMPLE IS DONE FOR ONLY A SINGLE TEST IMAGE FROM ds_train.batch(
for example in ds_train.batch(1).take(1):
    # Starts Profile logging
    tf.profiler.experimental.start(os.path.join(log_dir, f'single-{datetime.d

    # Actual inference
    # TODO: Your Code Here
    model.predict(example['image'])

    # Stops Profile logging
    tf.profiler.experimental.stop()

# THE BELOW EXAMPLE IS DONE FOR OUR THREE SAMPLE IMAGES COMING FROM sample_im
# tf.profiler.experimental.start(os.path.join(log_dir, f'single-{datetime.dat
# for image in sample_imgs:
#     # Starts Profile logging
#     #tf.profiler.experimental.start(os.path.join(log_dir, f'single-{datetim

#     #Actual inference
#     # TODO: Your Code Here
#     img = tf.keras.preprocessing.image.img_to_array(image['image'])
#     img = np.expand_dims(img, axis=0)
#     img = tf.keras.applications.vgg16.preprocess_input(img)

#     model.predict(img)

#     # Stops Profile logging
#     tf.profiler.experimental.stop()

# Load the TensorBoard notebook extension.
%load_ext tensorboard

# Launch TensorBoard and navigate to the Profile tab to view performance prof
# *** Please note just execute this command ones in a session and
# then logs for subsequent runs would be auto detected in tensorboard- url: h
%tensorboard --logdir=log_dir

# You could view the tensorboard in the browser url: http://localhost:6006/

```

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6006 (pid 226879), started 5 days, 20:47:26 ago.  
(Use '!kill 226879' to kill it.)



In [ ]:

```
# TODO: Sample Profiling - Online Inference:

# Vary this from 10, 100, 1000 to simulate multiple online inference
loop_index = [10, 100, 1000]

for idx in loop_index:
    # Starts Profile logging
    # tf.profiler.experimental.stop()
    tf.profiler.experimental.start(os.path.join(log_dir, f'online-inference-{i

    # Actual online inference
    # TODO: Your Code Here
    for index, example in enumerate(ds_val.take(idx)):
        sample = tf.keras.preprocessing.image.img_to_array(example['image'])
        sample = np.expand_dims(sample, axis=0)
        sample = tf.keras.applications.vgg16.preprocess_input(sample)

        model.predict(sample)
    print(f'BATCH CALCULATION {idx} COMPLETE')

    # Stops Profile logging
    tf.profiler.experimental.stop()

# Load the TensorBoard notebook extension.
%load_ext tensorboard

# Launch TensorBoard and navigate to the Profile tab to view performance prof
# *** Please note just execute this command ones in a session and
# then logs for subsequent runs would be auto detected in tensorboard- url: h
%tensorboard --logdir=log_dir

# You could view the tensorboard in the browser url: http://localhost:6006/ a
```

BATCH CALCULATION 10 COMPLETE

BATCH CALCULATION 100 COMPLETE

BATCH CALCULATION 1000 COMPLETE

The tensorboard extension is already loaded. To reload it, use:

%reload\_ext tensorboard

Reusing TensorBoard on port 6006 (pid 226879), started 5 days, 20:47:58 ago.

(Use '!kill 226879' to kill it.)



In [ ]:

```
# TODO: Sample Profiling - Batch Inference:

# We would only perform batch inference for a subset of validation set i.e. 1
# using different batch sizes of 20, 40, 100, 200

# Decides the size of the batch. Try: 20, 40, 100, 200
batch_size = [20, 40, 100, 200]

for batch in batch_size:
    # Starts Profile logging

    tf.profiler.experimental.start(os.path.join(log_dir, f'batch-{batch}-{dat

    # Actual Batch inference
    # TODO: Your Code Here
    for example in enumerate(ds_val.batch(batch).take(1000)):
        model.predict(sample)
    print(f'BATCH SIZE {batch} COMPLETE')

    # Stops Profile logging
    tf.profiler.experimental.stop()

# Load the TensorBoard notebook extension.
%load_ext tensorboard

# Launch TensorBoard and navigate to the Profile tab to view performance prof
# *** Please note just execute this command ones in a session and
# then logs for subsequent runs would be auto detected in tensorboard- url: h
%tensorboard --logdir=log_dir

# You could view the tensorboard in the browser url: http://localhost:6006/ a
```

BATCH SIZE 20 COMPLETE

BATCH SIZE 40 COMPLETE

BATCH SIZE 100 COMPLETE

BATCH SIZE 200 COMPLETE

The tensorboard extension is already loaded. To reload it, use:

%reload\_ext tensorboard

Reusing TensorBoard on port 6006 (pid 226879), started 5 days, 20:48:29 ago.

(Use '!kill 226879' to kill it.)

Training

In [ ]:

```
# Setup for model training
from tensorflow.keras import Model, datasets
from tensorflow.keras.models import Sequential
from tensorflow.keras.losses import categorical_crossentropy
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.layers import Dense, Flatten, Conv2D, AveragePooling2D,

train_dir = os.path.abspath(os.path.join('train_data', datetime.datetime.now(
pathlib.Path(train_dir).mkdir(exist_ok=True, parents=True))

# Using early stopping to monitor validation accuracy
callbacks = [
    tf.keras.callbacks.EarlyStopping(
        # Stop training when `val_loss` is no longer improving
        monitor="val_loss",
        # "no longer improving" being defined as "no better than 1e-2 less"
        min_delta=1e-4,
        # "no longer improving" being further defined as "for at least 2 epochs"
        patience=2,
        verbose=1,
    ),
    tf.keras.callbacks.TensorBoard(log_dir=train_dir, histogram_freq=1)
]
```

In [ ]:

```
# Basic CNN model
train_model = Sequential()

# conv1
train_model.add(Conv2D(32, (5, 5), input_shape=(64, 64, 3), activation='relu'))
train_model.add(Conv2D(32, (5,5),activation='relu'))
train_model.add(MaxPooling2D(pool_size=(2, 2)))
train_model.add(Conv2D(64, (3,3), activation='relu'))
train_model.add(Conv2D(64, (3,3), activation='relu'))
train_model.add(MaxPooling2D(pool_size=(2, 2)))
train_model.add(Conv2D(64, (3,3), activation='relu'))
train_model.add(Conv2D(128, (3,3), activation='relu'))
train_model.add(MaxPooling2D(pool_size=(2, 2)))
train_model.add(Flatten())

# fc1
train_model.add(Dense(256, activation='relu'))

# fc2
train_model.add(Dense(200, activation='softmax'))

train_model.compile(loss='categorical_crossentropy', optimizer=tf.keras.optim
train_model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_12 (Conv2D)	(None, 60, 60, 32)	2432
conv2d_13 (Conv2D)	(None, 56, 56, 32)	25632



max_pooling2d_6 (MaxPooling2)	(None, 28, 28, 32)	0
conv2d_14 (Conv2D)	(None, 26, 26, 64)	18496
conv2d_15 (Conv2D)	(None, 24, 24, 64)	36928
max_pooling2d_7 (MaxPooling2)	(None, 12, 12, 64)	0
conv2d_16 (Conv2D)	(None, 10, 10, 64)	36928
conv2d_17 (Conv2D)	(None, 8, 8, 128)	73856
max_pooling2d_8 (MaxPooling2)	(None, 4, 4, 128)	0
flatten_2 (Flatten)	(None, 2048)	0
dense_4 (Dense)	(None, 256)	524544
dense_5 (Dense)	(None, 200)	51400
=====		
Total params: 770,216		
Trainable params: 770,216		
Non-trainable params: 0		

In [ ]:

```

# TODO: Attempt to train your own model with different batch sizes

def normalize_img(image, label):
    return tf.cast(image, tf.float32) / 255., label

def to_categorical(image, label):
    label = tf.one_hot(tf.cast(label, tf.int32), 200)
    return tf.cast(image, tf.float32), tf.cast(label, tf.int64)

ds_re = tiny_imagenet_builder.as_dataset(as_supervised=True)
ds_retrain, ds_reval = ds_re["train"], ds_re["validation"]

ds_retrain = ds_retrain.cache().shuffle(1024)
ds_reval = ds_reval.cache().shuffle(1024)

ds_retrain = ds_retrain.map(normalize_img, num_parallel_calls=tf.data.AUTOTUNE)
ds_reval = ds_reval.map(normalize_img, num_parallel_calls=tf.data.AUTOTUNE)

ds_retrain = ds_retrain.map(to_categorical, num_parallel_calls=tf.data.AUTOTUNE)
ds_reval = ds_reval.map(to_categorical, num_parallel_calls=tf.data.AUTOTUNE)

epoch_size = 20

weights = train_model.get_weights()

for batch_size in [32, 64, 128]:
    # Setup our batched datasets
    # TODO: Your Code Here

    print(f"BATCH SIZE {batch_size}:")
    retrain = ds_retrain.batch(batch_size, num_parallel_calls=tf.data.AUTOTUNE)
    reval = ds_reval.batch(batch_size, num_parallel_calls=tf.data.AUTOTUNE)

    # Run training
    # TODO: Your Code Here
    train_model.set_weights(weights=weights)

    train_model.fit(x=retrain, validation_data=reval, epochs=epoch_size, call

    # Save the cnn model
    train_model.save(os.path.join(log_dir, f'CNN_TinyImageNet_train_batch{bat

    # TODO: Get the top-1 and top-5 of your newly trained model
    # TODO: Your Code Here

    total = acc_top1 = acc_top5 = 0

    # TODO: Your Code Here
    for image in ds_val.batch(batch_size):
        prediction = train_model.predict(image['image'].numpy())

        top_1 = tf.math.top_k(prediction, k=1).indices
        top_5 = tf.math.top_k(prediction, k=5).indices

        for index, labels in enumerate(image['label'].numpy()):
            if labels in top_1[index]:

```

```

        acc_top1 += 1
    if labels in top_5[index]:
        acc_top5 += 1

    total += batch_size

    print(f"Top 1: {(acc_top1 / total)* 100:.2f}%, Top 5: {(acc_top5 / total)

```

BATCH SIZE 32:

Epoch 1/20

3125/3125 [=====] - 18s 6ms/step - loss: 4.9877 - accuracy: 0.0294 - val\_loss: 4.6644 - val\_accuracy: 0.0615

Epoch 2/20

3125/3125 [=====] - 17s 5ms/step - loss: 4.4464 - accuracy: 0.0891 - val\_loss: 4.3280 - val\_accuracy: 0.1081

Epoch 3/20

3125/3125 [=====] - 16s 5ms/step - loss: 4.1393 - accuracy: 0.1311 - val\_loss: 4.0306 - val\_accuracy: 0.1478

Epoch 4/20

3125/3125 [=====] - 17s 5ms/step - loss: 3.8915 - accuracy: 0.1685 - val\_loss: 3.8682 - val\_accuracy: 0.1715

Epoch 5/20

3125/3125 [=====] - 16s 5ms/step - loss: 3.6877 - accuracy: 0.1999 - val\_loss: 3.7318 - val\_accuracy: 0.1915

Epoch 6/20

3125/3125 [=====] - 16s 5ms/step - loss: 3.5176 - accuracy: 0.2264 - val\_loss: 3.6333 - val\_accuracy: 0.2064

Epoch 7/20

3125/3125 [=====] - 16s 5ms/step - loss: 3.3771 - accuracy: 0.2500 - val\_loss: 3.5579 - val\_accuracy: 0.2213

Epoch 8/20

3125/3125 [=====] - 16s 5ms/step - loss: 3.2516 - accuracy: 0.2700 - val\_loss: 3.5205 - val\_accuracy: 0.2274

Epoch 9/20

3125/3125 [=====] - 16s 5ms/step - loss: 3.1331 - accuracy: 0.2905 - val\_loss: 3.4480 - val\_accuracy: 0.2410

Epoch 10/20

3125/3125 [=====] - 16s 5ms/step - loss: 3.0244 - accuracy: 0.3100 - val\_loss: 3.4701 - val\_accuracy: 0.2443

Epoch 11/20

3125/3125 [=====] - 16s 5ms/step - loss: 2.9181 - accuracy: 0.3292 - val\_loss: 3.4718 - val\_accuracy: 0.2477

Epoch 00011: early stopping

Top 1: 12.57%, Top 5: 18.51%

BATCH SIZE 64:

Epoch 1/20

1563/1563 [=====] - 12s 8ms/step - loss: 5.2673 - accuracy: 0.0068 - val\_loss: 5.1256 - val\_accuracy: 0.0142

Epoch 2/20

1563/1563 [=====] - 12s 7ms/step - loss: 5.0439 - accuracy: 0.0237 - val\_loss: 4.9358 - val\_accuracy: 0.0363

Epoch 3/20

1563/1563 [=====] - 11s 7ms/step - loss: 4.8173 - accuracy: 0.0504 - val\_loss: 4.7299 - val\_accuracy: 0.0615

Epoch 4/20

1563/1563 [=====] - 12s 7ms/step - loss: 4.5471 - accuracy: 0.0832 - val\_loss: 4.4268 - val\_accuracy: 0.0997

Epoch 5/20

1563/1563 [=====] - 11s 7ms/step - loss: 4.2747 - ac

```
curacy: 0.1167 - val_loss: 4.1854 - val_accuracy: 0.1235
Epoch 6/20
1563/1563 [=====] - 12s 7ms/step - loss: 4.0641 - ac
curacy: 0.1456 - val_loss: 4.0385 - val_accuracy: 0.1456
Epoch 7/20
1563/1563 [=====] - 12s 7ms/step - loss: 3.8868 - ac
curacy: 0.1724 - val_loss: 3.9180 - val_accuracy: 0.1645
Epoch 8/20
1563/1563 [=====] - 12s 7ms/step - loss: 3.7308 - ac
curacy: 0.1966 - val_loss: 3.8403 - val_accuracy: 0.1770
Epoch 9/20
1563/1563 [=====] - 12s 8ms/step - loss: 3.5929 - ac
curacy: 0.2165 - val_loss: 3.6954 - val_accuracy: 0.2025
Epoch 10/20
1563/1563 [=====] - 12s 7ms/step - loss: 3.4696 - ac
curacy: 0.2393 - val_loss: 3.6662 - val_accuracy: 0.2073
Epoch 11/20
1563/1563 [=====] - 11s 7ms/step - loss: 3.3604 - ac
curacy: 0.2579 - val_loss: 3.6044 - val_accuracy: 0.2197
Epoch 12/20
1563/1563 [=====] - 12s 7ms/step - loss: 3.2565 - ac
curacy: 0.2756 - val_loss: 3.5643 - val_accuracy: 0.2205
Epoch 13/20
1563/1563 [=====] - 11s 7ms/step - loss: 3.1687 - ac
curacy: 0.2913 - val_loss: 3.5538 - val_accuracy: 0.2293
Epoch 14/20
1563/1563 [=====] - 11s 7ms/step - loss: 3.0743 - ac
curacy: 0.3067 - val_loss: 3.5116 - val_accuracy: 0.2359
Epoch 15/20
1563/1563 [=====] - 12s 7ms/step - loss: 2.9883 - ac
curacy: 0.3237 - val_loss: 3.5024 - val_accuracy: 0.2398
Epoch 16/20
1563/1563 [=====] - 12s 7ms/step - loss: 2.8992 - ac
curacy: 0.3388 - val_loss: 3.5197 - val_accuracy: 0.2418
Epoch 17/20
1563/1563 [=====] - 12s 7ms/step - loss: 2.8182 - ac
curacy: 0.3521 - val_loss: 3.5207 - val_accuracy: 0.2461
Epoch 00017: early stopping
Top 1: 15.63%, Top 5: 22.04%
BATCH SIZE 128:
Epoch 1/20
782/782 [=====] - 10s 12ms/step - loss: 5.2967 - acc
uracy: 0.0053 - val_loss: 5.2869 - val_accuracy: 0.0051
Epoch 2/20
782/782 [=====] - 9s 12ms/step - loss: 5.1960 - accu
racy: 0.0101 - val_loss: 5.1091 - val_accuracy: 0.0124
Epoch 3/20
782/782 [=====] - 9s 12ms/step - loss: 5.0709 - accu
racy: 0.0205 - val_loss: 5.0285 - val_accuracy: 0.0228
Epoch 4/20
782/782 [=====] - 9s 12ms/step - loss: 4.9648 - accu
racy: 0.0328 - val_loss: 4.8737 - val_accuracy: 0.0418
Epoch 5/20
782/782 [=====] - 9s 12ms/step - loss: 4.7869 - accu
racy: 0.0546 - val_loss: 4.7144 - val_accuracy: 0.0648
Epoch 6/20
782/782 [=====] - 9s 12ms/step - loss: 4.5699 - accu
racy: 0.0800 - val_loss: 4.4717 - val_accuracy: 0.0930
Epoch 7/20
782/782 [=====] - 9s 12ms/step - loss: 4.3744 - accu
```

```
racy: 0.1034 - val_loss: 4.3184 - val_accuracy: 0.1139
Epoch 8/20
782/782 [=====] - 9s 12ms/step - loss: 4.2085 - accu
racy: 0.1270 - val_loss: 4.1920 - val_accuracy: 0.1268
Epoch 9/20
782/782 [=====] - 9s 12ms/step - loss: 4.0655 - accu
racy: 0.1463 - val_loss: 4.1014 - val_accuracy: 0.1444
Epoch 10/20
782/782 [=====] - 9s 12ms/step - loss: 3.9341 - accu
racy: 0.1668 - val_loss: 4.0131 - val_accuracy: 0.1532
Epoch 11/20
782/782 [=====] - 9s 12ms/step - loss: 3.8194 - accu
racy: 0.1855 - val_loss: 3.9124 - val_accuracy: 0.1684
Epoch 12/20
782/782 [=====] - 9s 12ms/step - loss: 3.7083 - accu
racy: 0.2013 - val_loss: 3.8172 - val_accuracy: 0.1824
Epoch 13/20
782/782 [=====] - 9s 12ms/step - loss: 3.6082 - accu
racy: 0.2195 - val_loss: 3.7911 - val_accuracy: 0.1848
Epoch 14/20
782/782 [=====] - 9s 12ms/step - loss: 3.5176 - accu
racy: 0.2328 - val_loss: 3.7242 - val_accuracy: 0.1980
Epoch 15/20
782/782 [=====] - 9s 12ms/step - loss: 3.4331 - accu
racy: 0.2462 - val_loss: 3.6923 - val_accuracy: 0.2025
Epoch 16/20
782/782 [=====] - 9s 12ms/step - loss: 3.3585 - accu
racy: 0.2585 - val_loss: 3.6665 - val_accuracy: 0.2112
Epoch 17/20
782/782 [=====] - 9s 12ms/step - loss: 3.2818 - accu
racy: 0.2722 - val_loss: 3.6654 - val_accuracy: 0.2122
Epoch 18/20
782/782 [=====] - 9s 12ms/step - loss: 3.2148 - accu
racy: 0.2823 - val_loss: 3.6430 - val_accuracy: 0.2138
Epoch 19/20
782/782 [=====] - 9s 12ms/step - loss: 3.1446 - accu
racy: 0.2940 - val_loss: 3.5986 - val_accuracy: 0.2234
Epoch 20/20
782/782 [=====] - 9s 12ms/step - loss: 3.0798 - accu
racy: 0.3053 - val_loss: 3.6327 - val_accuracy: 0.2232
```

In [ ]:

```

# TODO: Train your model with 3 different numbers of epochs
batch_size = 32

# Setup your datasets
# TODO: Your Code Here
retrain = ds_retrain.batch(batch_size, num_parallel_calls=tf.data.AUTOTUNE)
reval = ds_reval.batch(batch_size, num_parallel_calls=tf.data.AUTOTUNE)

for epoch_size in [3, 10, 100]:
    # Run training
    # TODO: Your Code Here
    train_model.set_weights(weights=weights)
    train_model.fit(x=retrain, validation_data=reval, epochs=epoch_size, call

    # Save the cnn model
    train_model.save(os.path.join(log_dir, f'CNN_TinyImageNet_train_epoch{epoch_size}.h5'))

    # TODO: Get the top-1 and top-5 of your newly trained model
    total = acc_top1 = acc_top5 = 0

    # TODO: Your Code Here
    for image in ds_val.batch(batch_size):
        prediction = train_model.predict(image['image'].numpy())

        top_1 = tf.math.top_k(prediction, k=1).indices
        top_5 = tf.math.top_k(prediction, k=5).indices

        for index, labels in enumerate(image['label'].numpy()):
            if labels in top_1[index]:
                acc_top1 += 1
            if labels in top_5[index]:
                acc_top5 += 1

        total += batch_size

    print(f"Top 1: {(acc_top1 / total)* 100:.2f}%, Top 5: {(acc_top5 / total)* 100:.2f}%")

```

Epoch 1/3

3125/3125 [=====] - 17s 5ms/step - loss: 5.2374 - accuracy: 0.0088 - val\_loss: 5.0587 - val\_accuracy: 0.0197

Epoch 2/3

3125/3125 [=====] - 16s 5ms/step - loss: 4.8771 - accuracy: 0.0401 - val\_loss: 4.6827 - val\_accuracy: 0.0629

Epoch 3/3

3125/3125 [=====] - 16s 5ms/step - loss: 4.4578 - accuracy: 0.0912 - val\_loss: 4.2967 - val\_accuracy: 0.1106

Top 1: 8.11%, Top 5: 14.42%

Epoch 1/10

3125/3125 [=====] - 17s 5ms/step - loss: 5.1722 - accuracy: 0.0146 - val\_loss: 4.9507 - val\_accuracy: 0.0316

Epoch 2/10

3125/3125 [=====] - 16s 5ms/step - loss: 4.7465 - accuracy: 0.0555 - val\_loss: 4.5753 - val\_accuracy: 0.0793

Epoch 3/10

3125/3125 [=====] - 16s 5ms/step - loss: 4.3093 - accuracy: 0.1093 - val\_loss: 4.1985 - val\_accuracy: 0.1167

Epoch 4/10

3125/3125 [=====] - 17s 5ms/step - loss: 4.0029 - accuracy: 0.1442 - val\_loss: 3.9507 - val\_accuracy: 0.1316

```
curacy: 0.1512 - val_loss: 3.9361 - val_accuracy: 0.1601
Epoch 5/10
3125/3125 [=====] - 17s 5ms/step - loss: 3.7801 - ac
curacy: 0.1838 - val_loss: 3.7881 - val_accuracy: 0.1826
Epoch 6/10
3125/3125 [=====] - 16s 5ms/step - loss: 3.6053 - ac
curacy: 0.2126 - val_loss: 3.6934 - val_accuracy: 0.1981
Epoch 7/10
3125/3125 [=====] - 17s 5ms/step - loss: 3.4610 - ac
curacy: 0.2356 - val_loss: 3.5876 - val_accuracy: 0.2141
Epoch 8/10
3125/3125 [=====] - 16s 5ms/step - loss: 3.3265 - ac
curacy: 0.2578 - val_loss: 3.5800 - val_accuracy: 0.2192
Epoch 9/10
3125/3125 [=====] - 16s 5ms/step - loss: 3.2068 - ac
curacy: 0.2798 - val_loss: 3.5652 - val_accuracy: 0.2233
Epoch 10/10
3125/3125 [=====] - 16s 5ms/step - loss: 3.0987 - ac
curacy: 0.2981 - val_loss: 3.5036 - val_accuracy: 0.2315
Top 1: 11.56%, Top 5: 18.27%
Epoch 1/100
3125/3125 [=====] - 17s 5ms/step - loss: 5.1788 - ac
curacy: 0.0130 - val_loss: 5.0259 - val_accuracy: 0.0238
Epoch 2/100
3125/3125 [=====] - 16s 5ms/step - loss: 4.8665 - ac
curacy: 0.0434 - val_loss: 4.7043 - val_accuracy: 0.0642
Epoch 3/100
3125/3125 [=====] - 17s 5ms/step - loss: 4.4663 - ac
curacy: 0.0912 - val_loss: 4.2927 - val_accuracy: 0.1127
Epoch 4/100
3125/3125 [=====] - 16s 5ms/step - loss: 4.1179 - ac
curacy: 0.1381 - val_loss: 4.0086 - val_accuracy: 0.1508
Epoch 5/100
3125/3125 [=====] - 16s 5ms/step - loss: 3.8488 - ac
curacy: 0.1769 - val_loss: 3.8156 - val_accuracy: 0.1808
Epoch 6/100
3125/3125 [=====] - 17s 5ms/step - loss: 3.6498 - ac
curacy: 0.2077 - val_loss: 3.6929 - val_accuracy: 0.1982
Epoch 7/100
3125/3125 [=====] - 17s 5ms/step - loss: 3.4875 - ac
curacy: 0.2334 - val_loss: 3.6326 - val_accuracy: 0.2098
Epoch 8/100
3125/3125 [=====] - 17s 5ms/step - loss: 3.3483 - ac
curacy: 0.2579 - val_loss: 3.5373 - val_accuracy: 0.2267
Epoch 9/100
3125/3125 [=====] - 17s 5ms/step - loss: 3.2238 - ac
curacy: 0.2782 - val_loss: 3.5222 - val_accuracy: 0.2310
Epoch 10/100
3125/3125 [=====] - 16s 5ms/step - loss: 3.1063 - ac
curacy: 0.2990 - val_loss: 3.4915 - val_accuracy: 0.2378
Epoch 11/100
3125/3125 [=====] - 17s 5ms/step - loss: 2.9929 - ac
curacy: 0.3198 - val_loss: 3.4371 - val_accuracy: 0.2469
Epoch 12/100
3125/3125 [=====] - 16s 5ms/step - loss: 2.8854 - ac
curacy: 0.3391 - val_loss: 3.4653 - val_accuracy: 0.2505
Epoch 13/100
3125/3125 [=====] - 17s 5ms/step - loss: 2.7797 - ac
curacy: 0.3568 - val_loss: 3.4933 - val_accuracy: 0.2475
Epoch 00013: early stopping
```



Top 1: 14.82%, Top 5: 21.20%

## Above and Beyond

```
In [ ]: # Benchmark our dataset to make sure loading our data isn't a bottleneck ...  
# (This can be skipped since it can take a bit and isn't all that important)  
  
tfds.benchmark(ds_train.batch(32), batch_size=32, num_iter=2**20)
```

```
WARNING:absl:Number of iteration shorter than expected (3123 vs 1048576)  
***** Summary *****
```

```
Examples/sec (First included) 140740.91 ex/sec (total: 100000 ex, 0.71 sec)  
Examples/sec (First only) 2236.96 ex/sec (total: 32 ex, 0.01 sec)  
Examples/sec (First excluded) 143586.72 ex/sec (total: 99968 ex, 0.70 sec)
```

Out[ ]: **BenchmarkResult:**

	duration	num_examples	avg
<b>first+lasts</b>	0.710525	100000	140740.911071
<b>first</b>	0.014305	32	2236.964900
<b>lasts</b>	0.696220	99968	143586.723272

In [ ]:

```

# RUN INFERENCE ON A CUSTOM IMAGE AND DISPLAY FEATURE MAPS OF THE CAMPANILE!

# We can again view the layer outputs as well. Here we are pretending they are
# TODO: Visualize the 2 convolutional layers outputs (intermediate feature maps)

# TODO: Your Code Here

# Run inference to get feature maps on earlier layer
sub_model = Model(inputs=model.inputs, outputs=model.layers[0].output)

image = tf.keras.preprocessing.image.load_img('campanile.jpg', target_size=(64, 64))
image = tf.keras.preprocessing.image.img_to_array(image)
image = np.expand_dims(image, axis=0)
image = tf.keras.applications.vgg16.preprocess_input(image)

feature_maps = sub_model.predict(image)
idx = 1
fig = plt.figure(figsize=(20, 20))
for i in range(1, feature_maps.shape[3]+1):

    if idx == 65:
        break

    plt.subplot(8, 8, i)
    plt.imshow(feature_maps[0, :, :, i-1])
    idx += 1

plt.show()

print(f'Beginning Layer: {model.layers[0].name}') # Prove name of layer
print(f'Data Type: {feature_maps.dtype}') # Print dtype
print(f'Dimensions: {feature_maps.ndim} total {feature_maps.shape}') # Print shape
print(f'Required memory: {feature_maps.nbytes}B ({feature_maps.nbytes/1000}KB)

# Run inference to get feature maps on end conv layer

sub_model = Model(inputs=model.input, outputs=model.layers[7].output)

#print(f"Printing feature maps for layer {model.layers[7].name}...")

image = sample_imgs[0]['image']
image = tf.keras.preprocessing.image.img_to_array(image)
image = np.expand_dims(image, axis=0)
image = tf.keras.applications.vgg16.preprocess_input(image)

feature_maps = sub_model.predict(image)

idx = 1
fig = plt.figure(figsize=(20, 20))
for i in range(1, feature_maps.shape[3]+1):

    if idx == 65:
        break


    plt.subplot(8, 8, i)
    plt.imshow(feature_maps[0, :, :, i-1])

```

```
idx += 1

plt.show()

print(f'End Layer: {model.layers[7].name}') # Prove name of layer
print(f'Data Type: {feature_maps.dtype}') # Print dtype
print(f'Dimensions: {feature_maps.ndim} total {feature_maps.shape}') # Print
print(f'Required memory: {feature_maps.nbytes}B ({feature_maps.nbytes/1000}KB
```



Beginning Layer: conv2d  
Data Type: float32  
Dimensions: 4 total (1, 60, 60, 32)  
Required memory: 460800B (460.8KB)

